

Binaml: Continual Learning over Binary Feature Streams

Romain Zimmer
hi@romainzimmer.com

2026-08-11 (Draft)

Abstract

Binaml focuses on continual learning models over streams of binary features subject to data drift. Binary features provide a task-agnostic representation and enable models optimized for memory, latency, and energy efficiency. This paper specifies Binaml’s current regression model, **BRegressor**, as an ensemble of batch-learned boolean functions with scalar weights rather than a monolithic growing feature store. It adapts online rather than relying on a stationary train/serve assumption and learns compact compositional representations from residual feedback. The model is environment-agnostic: it does not depend on how the stream is produced. This paper specifies the current regression model and reports reproducible prequential comparisons on versioned synthetic streams.

1 Notation

d, x_t, y_t Input width, binary input, and scalar regression target.

$\hat{y}_t, e_t, s_t$ Prediction, residual, and residual-sign label.

b, w_k, η, λ Intercept, ensemble weight of function k , learning rate, and ℓ_2 decay.

$f_k(x)$ Stored boolean function k mapping $\{0, 1\}^d$ to $\{0, 1\}$, stored as a compact boolean function.

$T \in \{0, \dots, 15\}$ Four-bit truth table of an arity-2 composed node.

$B, S, n, K, L_{\max}, K_{\max}$ Feature batch size, SGD steps per observation, batch index, parent top- K width, maximum builder depth, and ensemble capacity.

2 Model

At time t , the model receives $x_t \in \{0, 1\}^d$ and later observes a finite scalar y_t . If useful signal is sparse weighted boolean structure over bits, a linear head over learned boolean indicators is a direct hypothesis class: continuous amplitudes, discrete features, and an inspectable composition. Let K_t be the number of stored functions after batch finalizations through time t . The prediction is

$$\hat{y}_t = b + \sum_{k=1}^{K_t} w_k (2f_k(x_t) - 1),$$

where $b \in \mathbb{R}$, each $w_k \in \mathbb{R}$, and each $f_k(x_t) \in \{0, 1\}$ is centered before entering the linear head:

$$2f_k(x_t) - 1 \in \{-1, 1\}.$$

Each f_k is a compact boolean function built from one residual-sign batch, not a node in a shared graph.

Composed features. Composition starts from the most basic nontrivial boolean learning unit: two-input tables. Combining them yields richer features without beginning from high-arity search. A composed node combines two parent activations using a four-bit truth table $T \in \{0, \dots, 15\}$.

3 Feature representation and batch learning

Each feature batch grows an ephemeral wide DAG: layer zero holds the d source bits; at each subsequent layer the builder forms every distinct pair among the top- K parent nodes by in-sample match score, learns a truth table on the current batch, and appends the resulting node. Depth never exceeds $L_{\max}$. When the batch fills, the builder selects one output node by descending in-sample accuracy, then depth, then node id, compacts the subgraph into a **FunctionGraph**, and appends it to the ensemble with weight zero. If the ensemble exceeds capacity $K_{\max}$, the function with smallest $|w_k|$ is removed, breaking ties toward the oldest index. There is no hold-out batch and no shared candidate queue across batches.

3.1 Compact evaluation

A stored **FunctionGraph** lists source indices and composed nodes in topological order. Each node is either a source coordinate, a constant, or a composed operation referencing two already evaluated positions and one four-bit truth table. One forward pass evaluates $f_k(x)$ for prediction and SGD. The eight u8 counters used to learn a table cover every $(u, v, s) \in \{0, 1\}^3$ combination, which limits the feature batch size to $B \leq 255$.

Type	Fields or representation	Role
FunctionBuildConfig	{batch_size, parent_top_k, max_layers}	batch-local graph growth limits
FunctionBuilder	stateless builder	grows ephemeral DAG from one SignBatch
FunctionGraph	sources, compact nodes, output index	immutable stored function
CompactNode	Source, Constant, or Composed{first, second, truth_table}	one evaluation step
BRegressor	functions: Vec<FunctionGraph>, weights: Vec<f64>, intercept, batch buffers, count	online ensemble state

4 Online linear-head updates

Function identity is discrete; weights are continuous, so the two are updated separately. For each observation (x_t, y_t) , evaluate the current stored functions once, compute the residual sign with the current weights, then perform S single-sample squared-error gradient steps on that activation row. The newly appended function at a batch boundary is excluded from SGD on the batch-fill step because it is added only after those updates. For one such step with pre-step residual $e_t = y_t - \hat{y}_t$, learning rate $\eta > 0$, and decay $\lambda \geq 0$, apply

$$\begin{aligned}
 w_k &\leftarrow (1 - \eta\lambda) w_k && \text{for all stored } k, \\
 b &\leftarrow b + \eta e_t, \\
 w_k &\leftarrow w_k + \eta e_t (2f_k(x_t) - 1).
 \end{aligned}$$

There is no gradient through boolean structure: SGD tracks residual amplitude at the sample level, while structure search runs when enough residual-sign evidence exists. The residual sign

$$s_t = \mathbf{1}\{e_t \geq 0\}$$

is computed with the current weights before the associated linear updates and appended to the current feature batch. After every B new observations, that batch is consumed to grow, select, compact, and append one function, then cleared.

5 Residual-sign supervision and batching

Feature learning targets residual signs, not y_t itself. Residual signs are a binary supervision bit for the discrete builder, while the linear head absorbs magnitude. Exact truth-table counting needs a bounded window; batching amortizes graph growth. Batch n is

$$\mathcal{B}_n = \{(x_t, s_t)\}_{t \in I_n}, \quad |I_n| = B,$$

where I_n is the contiguous block of B observations since the previous structural update. Processing $\mathcal{B}_n$ grows an ephemeral graph, selects one output by in-sample accuracy, compacts it, appends $(f, w = 0)$ to the ensemble, and prunes to $K_{\max}$ if needed. There is no hold-out batch.

Match score. For a boolean stream $z_{1:B}$ and residual signs $s_{1:B}$,

$$\sigma = \sum_{i=1}^B \mathbf{1}\{z_i = s_i\}.$$

Sources use their coordinate columns; composed nodes use their evaluated bitstreams.

6 Truth-table learning

Given two parent bitstreams $u_{1:B}$, $v_{1:B}$ and residual signs $s_{1:B}$, the learner returns the Hamming-error-minimizing four-bit table. Its state is a fixed array of eight `u8` counters, one for each combination of the two parent values and residual sign. The learned result is a four-bit truth table represented by one `u8`. The counters record the frequencies needed to select each table output by majority, with ties resolving to zero. As one counter may receive every observation, the `u8` counters bound the admissible batch size to $B \leq 255$.

7 Batch graph growth

For each layer $\ell = 1, \dots, L_{\max}$:

1. Let P be the top- K nodes in layer $\ell - 1$ by in-sample match score, breaking ties by node id.
2. Form every distinct pair in P .
3. For each pair, learn a truth table on $\mathcal{B}_n$ and append the resulting composed node to layer ℓ .

If P has fewer than two nodes, layer ℓ is skipped. After all layers, the builder selects the output node by descending match score, then depth, then id.

8 Ensemble pruning

When appending a new function would exceed $K_{\max}$, remove the stored index with smallest $|w_k|$, breaking ties toward the oldest index. Zero-weight experts can therefore rotate at capacity even when still structurally useful on past batches.

9 Online learning procedure

Configuration

Linear Learning rate η , decay λ , and S SGD steps per observation.

Batch structure Feature batch size B , parent top- K width K , maximum builder depth $L_{\max}$, and ensemble capacity $K_{\max}$.

Initialize: empty ensemble, $b \leftarrow 0$, and empty feature-batch buffers.

For each observation (x_t, y_t) :

1. Validate the input width and finite target.
2. Evaluate stored $f_k(x_t)$ and predict $\hat{y}_t = b + \sum_k w_k(2f_k(x_t) - 1)$.
3. Compute $s_t = \mathbf{1}\{y_t - \hat{y}_t \geq 0\}$ and append (x_t, s_t) to the current feature batch.
4. Apply S single-sample linear updates using the current activation row.
5. After B accumulated observations, grow a graph on the batch, select and compact one output, append it with weight zero, prune to $K_{\max}$ if needed, and clear the feature batch.

Cost. Prediction and activation evaluation are linear in the number of stored functions and their node counts. Each linear step costs one function count. Structural work occurs once per feature batch: graph growth scans selected parent pairs over B rows, then compaction produces one stored function.

10 API and integration boundaries

The `binaml.models` package owns the online model and has no environment dependency: the goal is a generic model, decoupled from where the data comes from and how it was produced. The public surface is `BRegressor`, which implements `predict(features)` and `observe(features, target)` for binary feature vectors of width d . The prediction call may retain pairing state; `observe` must follow a preceding `predict`. Feature learning is internal to `observe` once a batch fills. Evaluation packages compose models with environments through predict-then-observe prequential evaluation. Named benchmarks only compose these independent components.

11 Prequential evaluation protocol

For each emitted (x_t, y_t) , an evaluator first records the model prediction, then exposes the target through `observe`. Mean squared error is computed from the recorded pre-update predictions. Structure updates occur only after a full feature batch fills inside `observe`; residual signs are measured with the current weights before the corresponding linear updates. Hyperparameters, implementation version, and the online protocol are required to interpret any reported trajectory.

12 Evaluation

The benchmark inputs and model configuration are versioned in this paper’s benchmark directory. The task uses 1,000 observations and seeds $\{0, 1, 2, 3, 4\}$ with 32 input bits, 12 fixed target components, noise standard deviation 0.1, and the same function, DAG, and target-scale settings. Input, gate, and intercept distributions each resample independently with probability 0.01.

All models receive the scenario width. **BRegressor** uses learning rate $\eta = 5 \times 10^{-3}$, ℓ_2 decay 10^{-4} , feature batch size $B = 16$, five single-sample SGD steps, parent top- $K = 8$, builder depth $L_{\max} = 3$, and ensemble capacity $K_{\max} = 64$. Each stored function enters the linear head as $2f_k(x) - 1$. The linear baseline is a Python+NumPy replay-batch implementation using learning rate 0.03, the same decay, replay size 32, and step count, with uncentered binary inputs. The MLP baseline is a Scikit-learn wrapper using replay size 32 and three `partial_fit` passes; one 50-unit ReLU hidden layer; scikit-learn SGD with constant learning rate 0.003, $\alpha = 10^{-4}$, `power_t`=0.5, shuffling enabled, random state zero, tolerance 10^{-4} , zero momentum, no Nesterov momentum, and no early stopping. Its remaining configured values are validation fraction 0.1, $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$, 10 no-change iterations, and disabled verbose output; the wrapper fixes one iteration and disables warm start.

Tables report the mean $\pm$ standard error over five seeds. MSE is the optimized loss; RMSE expresses the same error in target units. Total time is the prequential predict-plus-observe time for a 1,000-sample trajectory.

Model	MSE	RMSE	total s
BRegressor (ours)	0.408 $\pm$ 0.018	0.639 $\pm$ 0.014	0.064 $\pm$ 0.012
Linear SGD (Python+NumPy)	0.745 $\pm$ 0.015	0.863 $\pm$ 0.008	0.172 $\pm$ 0.005
MLP SGD (Scikit-learn)	0.842 $\pm$ 0.024	0.917 $\pm$ 0.013	3.133 $\pm$ 0.201

13 Limitations and scope

This document specifies the ensemble regressor used by Binaml. It does not claim optimality of residual-sign learning or in-sample function selection. Ensemble churn at capacity can drop zero-weight experts that remain useful on past batches; in-sample output selection trades hold-out promotion for simpler batch boundaries. The reported comparison is limited to the versioned synthetic task, fixed configurations, and reproducible baseline runs described above.